

HOMEWORK 2

#Question 1: What are HDFS and YARN?

HDFS and YARN:

HDFS is the distributed file system in Hadoop for storing big data.

MapReduce is the processing framework for processing vast data in the Hadoop cluster in a distributed manner.

YARN is responsible for managing the resources amongst applications in the cluster.

#Question 2: What are the various Hadoop daemons and their roles in a Hadoop cluster?

a) various Hadoop daemons:

NameNode

DataNode

Secondary Name Node

Resource Manager

Node Manager

b) Role of the Hadoop daemons:

The Node Manager works on the Slaves System that manages the memory resource within the Node and Memory Disk.

Each Slave Node in a Hadoop cluster has a single NodeManager Daemon running in it. It also sends this monitoring information to the Resource Manager.

#Question 3: Why does one remove or add nodes in a Hadoop cluster frequently?

Basically, in a Hadoop cluster a Manager node will be deployed on a reliable hardware with high configurations, the Slave node's will be deployed on commodity hardware. So chance's of data node crashing is more .

So more frequently you will see admin's remove and add new data node's in a cluster.

#Question 4: What happens when two clients try to access the same file in the HDFS?
Multiple clients can't write into HDFS file at the similar time. When a client is granted a permission to write data on data node block, the block gets locked till the completion of a write operation. If some another client request to write on the same block of the same file then it is not permitted to do so.

#Question 5: How does NameNode tackle DataNode failures?
Data blocks on the failed Datanode are replicated on other Datanodes based on the specified replication factor in `hdfs-site.xml` file. Once the failed datanodes comes back the Name node will manage the replication factor again.
This is how Namenode handles the failure of data node.

#Question 6: What will you do when NameNode is down?
When the NameNode goes down, the file system goes offline. There is an optional SecondaryNameNode that can be hosted on a separate machine. It only creates checkpoints of the namespace by merging the edits file into the `fsimage` file and does not provide any real redundancy.

#Question 7: How is HDFS fault tolerant?
The HDFS is highly fault-tolerant that if any machine fails, the other machine containing the copy of that data automatically become active.
Distributed data storage -
This is one of the most important features of HDFS that makes Hadoop very powerful. Here, data is divided into multiple blocks and stored into nodes

#Question 8: Why do we use HDFS for applications having large data sets and not when there are a lot of small files?
HDFS is more efficient for a large number of data sets, maintained in a single file as compared to the small chunks of data stored in multiple files.

#Question 9:How do you define “block” in HDFS? What is the default block size in Hadoop 1 and in Hadoop 2? Can it be changed?

a)Blocks are the smallest continuous location on your hard drive where data is stored. HDFS stores each file as blocks, and distribute it across the Hadoop cluster.

b)The default size of a block in HDFS is 128 MB (Hadoop 2.x) and 64 MB (Hadoop 1.x)

c)Yes we can change the Hadoop block size.